Alexandra Reyes — MLOps Engineer

Years of Experience: 10
No Formal Degree in Statistics

Alexandra Reyes
MLOps Engineer
San Francisco, CA - (555) 123■4567 - alexandra.reyes@email.com - LinkedIn: linkedin.com/in/alexandra-reyes

SUMMARY

Data■driven professional with 10 years of experience in analytics, machine learning engineering, and infrastructure automation. Proven track record of building end■to■end ML pipelines, orchestrating containerized workloads, and implementing CI/CD processes at scale. Adept at collaborating across data science, software engineering, and operations teams to accelerate model delivery while ensuring robust, scalable, and highly available systems.

SKILLS

Python, SQL, Bash, Git, Docker, Kubernetes, Terraform, Jenkins, GitHub Actions, Apache Spark, Airflow, MLflow, TensorFlow, PyTorch, AWS, GCP, CI/CD, Prometheus, Grafana, Ansible, Tableau, Power BI

EXPERIENCE

- Built end■to■end MLOps pipelines integrating data ingestion, model training, validation, and deployment across multiple cloud environments.
- Managed model registry and versioning with MLflow and Amazon S3, ensuring reproducibility and traceability of every model artifact.
- Orchestrated data workflows in Airflow, enabling scheduled training and automated drift monitoring for production models.
- Implemented automated rollback and canary deployments for ML containers on Kubernetes, reducing downtime risk to <1 %.
- Partnered with product, data science, and SRE teams to define SLAs, monitor model health, and refine monitoring dashboards in Grafana.

- Designed and maintained CI/CD pipelines in Jenkins and GitHub Actions for microservice deployments, cutting release cycle time by 35 %.
- Provisioned and managed Kubernetes clusters on GKE, deploying Helm charts, configuring RBAC, and automating cluster scaling.
- Automated infrastructure provisioning with Terraform, reducing rollout time from days to minutes.
- Implemented Prometheus/Grafana monitoring and alerting for reliability metrics, achieving a 99.9 % uptime SLA.

- Designed and trained predictive models (Random Forest, Gradient Boosting, Neural Networks) for customer churn and credit■risk use cases.
- Developed feature■engineering pipelines with Spark and Pandas, improving model F1■score by 12 %.
- Deployed models to AWS SageMaker and GCP AI Platform, managing versioning and A/B testing to support continuous experimentation.
- Collaborated with data scientists to translate research prototypes into production■ready code, reducing model iteration latency by 40 %.

- Collected, cleaned, and normalized large datasets using SQL and Python to support strategic business decisions.
- Built interactive dashboards in Tableau and Power BI, increasing stakeholder data accessibility and reducing decision lag.
- Conducted statistical analyses and hypothesis testing to uncover trends and inform product roadmaps.
- Automated weekly report generation with cron jobs and Bash scripts, cutting manual effort by 45 %.

EDUCATION

Coursera (Online) - Certificate, Statistics | 2013

--