

Arjun Raghavan — MLOps Engineer

Years of Experience: 12
Associate in Statistics

ARJUN RAGHAVAN
MLOPS ENGINEER

Remote - (555) 123-4567 - arjun.raghavan@example.com - LinkedIn: linkedin.com/in/arjurn

SUMMARY

Accomplished MLOps Engineer with 12 years of experience in data analytics, machine learning, and infrastructure engineering across remote distributed teams. Proven track record in building scalable ML pipelines, containerizing models, automating CI/CD workflows, and implementing robust monitoring solutions on multi-cloud environments. Adept at bridging the gap between data science and production engineering to deliver high-impact AI solutions.

SKILLS

MLOps, Kubernetes, Docker, CI/CD, GitLab CI, Jenkins, Terraform, AWS, GCP, Azure, Python, TensorFlow, PyTorch, Scikit-learn, SQL, Linux, Bash, Airflow, Kubeflow, Prometheus, Grafana, S3, EKS, SageMaker, CloudWatch

EXPERIENCE

- Managed the full lifecycle of AI model deployments on AWS SageMaker and GCP Vertex AI.
- Automated model retraining pipelines triggered by data-drift detection, reducing failure rate.
- Integrated CI/CD pipelines with GitHub Actions, achieving zero-downtime deployments.
- Developed observability solutions using CloudWatch, ELK stack, and custom metrics to monitor model health and performance.
- Led a team of five ML engineers in designing scalable AI solutions for global clients.
- Implemented Infrastructure-as-Code with Terraform, provisioning resources across AWS, Azure, and GCP.
- Established monitoring and alerting dashboards using Prometheus and Grafana.
- Mentored junior staff on best practices for model versioning, reproducibility, and CI/CD.
- Architected end-to-end machine learning pipelines with Kubeflow and Docker.
- Containerized models and orchestrated deployments on Kubernetes across AWS EKS and GCP GKE.
- Developed CI/CD workflows with GitLab CI, ensuring rapid and reliable model rollouts.
- Optimized model inference performance, reducing latency by 40%.
- Built predictive models for customer churn and sales forecasting using XGBoost and Scikit-learn.
- Implemented feature engineering pipelines and performed hyperparameter tuning.
- Deployed models as RESTful services with Flask and Docker, enabling real-time scoring.
- Collaborated with product teams to translate model outputs into actionable business features.
- Collected, cleansed, and transformed large datasets using SQL and Python.
- Designed and maintained interactive dashboards for cross-functional stakeholders.
- Performed exploratory data analysis to identify trends and actionable insights.
- Automated data pipelines with Airflow, reducing report generation time by 30%.

EDUCATION

Associate in Statistics - Maplesoft Community College - Cedarville, USA - 2012