Dr. Maya Patel — Machine Learning Engineer

Years of Experience: 13
Ph.D. in Data Science

DR. MAYA PATEL
Machine Learning Engineer
Boston, MA - (555) 123-4567 - maya.patel@example.com - LinkedIn: linkedin.com/in/mayapatel

SUMMARY

Seasoned Machine Learning Engineer with 13 years of experience leading data■centric projects in academia and industry. Ph.D. from MIT where I pioneered scalable probabilistic models and advanced deep■learning architectures. Proven track record of designing production■grade ML pipelines, optimizing inference for low■latency applications, and mentoring cross■functional teams. Expertise spans supervised learning, NLP, computer vision, reinforcement learning, cloud MLOps, and distributed training.

- Machine Learning
- Deep Learning
- Natural Language Processing
- Computer Vision
- Reinforcement Learning
- Scalable Probabilistic Modeling
- Statistical Analysis
- Python
- PyTorch
- TensorFlow
- Keras
- scikit■learn
- SQL
- NoSQL
- Apache Spark
- Docker
- Kubernetes
- AWS SageMaker
- Google Cloud AI Platform
- Azure Machine Learning
- Git
- CI/CD
- MLOps
- Distributed Training
- Model Optimization
- TensorRT
- CUDA

EXPERIENCE

- Collected and cleaned heterogeneous retail data, improving data quality by 20 %.
- Developed ETL pipelines with Python and SQL, reducing processing time by 30 %.
- Performed exploratory data analysis and built Tableau dashboards, guiding marketing decisions that increased conversion rates by 12 %.

- Conducted research on scalable Bayesian inference for large■scale recommendation systems.
- Authored five peer■reviewed publications in top AI conferences.
- Supervised two undergraduate interns on data■engineering projects.
- Implemented prototype systems on Apache Spark that processed 1 TB of data weekly.
- Collaborated with product teams to validate models against production traffic.

- Led a research project on deep reinforcement learning for autonomous navigation.

- Developed and benchmarked state■of■the■art algorithms using PyTorch, achieving a 15 % improvement over baseline.
- Presented findings at NeurIPS and ICML conferences.
- Secured a $250,000 NSF research grant to fund ongoing work.

- Architected end■to■end recommendation pipelines serving 300 M daily users.
- Reduced inference latency by 35 % through model distillation and ONNX conversion.
- Built a scalable MLOps pipeline on AWS SageMaker, Docker, and Kubernetes, cutting deployment time from 48 hours to 3 hours.
- Mentored three junior engineers and organized internal workshops on best practices.
- Collaborated with cross■functional teams across product, data, and infrastructure to align ML models with business goals.

EDUCATION

Ph.D., Data Science - Massachusetts Institute of Technology, Cambridge, MA
2015

M.Sc., Computer Science - University of Illinois Urbana■Champaign, Urbana■Champaign, IL
2009

B.Sc., Computer Science - University of Illinois Urbana■Champaign, Urbana■Champaign, IL
2007