

Liam Chen — Machine Learning Engineer

Years of Experience: 7
Diploma in Data Science

LIAM CHEN

Machine Learning Engineer
San Francisco, CA - (555) 123-4567 - email@example.com - LinkedIn: linkedin.com/in/example

SUMMARY

Dedicated Machine Learning Engineer with 7 years of experience designing, deploying, and optimizing predictive models across diverse industries. Adept at translating complex data into actionable insights and leading cross-functional teams to deliver scalable AI solutions.

- Python, R, SQL
- TensorFlow, PyTorch, scikit-learn
- Apache Spark, Hadoop, PySpark
- AWS SageMaker, Docker, Kubernetes
- Data Visualization, Tableau, Pandas, NumPy

EXPERIENCE

- Headed a team of 4 ML engineers, delivering AI-driven product features that increased user engagement by 18%.
- Architected distributed training pipelines using Docker and Kubernetes, reducing training time by 35%.
- Implemented CI/CD workflows for ML models, cutting deployment lead time from days to hours.
- Designed end-to-end ML pipelines with PySpark, processing 10TB of streaming data nightly.
- Deployed models to AWS SageMaker, scaling effortlessly to support 50K concurrent inference requests.
- Led A/B testing frameworks that informed model selection and improved prediction accuracy by 12%.
- Assisted in building a collaborative filtering recommendation engine that lifted click-through rates by 9%.
- Conducted feature engineering and hyperparameter tuning, boosting model F1-score by 15%.
- Collaborated with data scientists to validate models and optimize training pipelines for production.
- Cleansed and transformed large datasets using SQL and Python, improving data quality for downstream analysis.
- Developed interactive Tableau dashboards, enabling real-time KPI monitoring for senior stakeholders.
- Built predictive models (linear regression, decision trees) to forecast monthly sales, achieving 85% accuracy.

EDUCATION

Diploma - Data Science
TechLaunch Academy, San Jose, CA
2017

--